

# Binary Classification

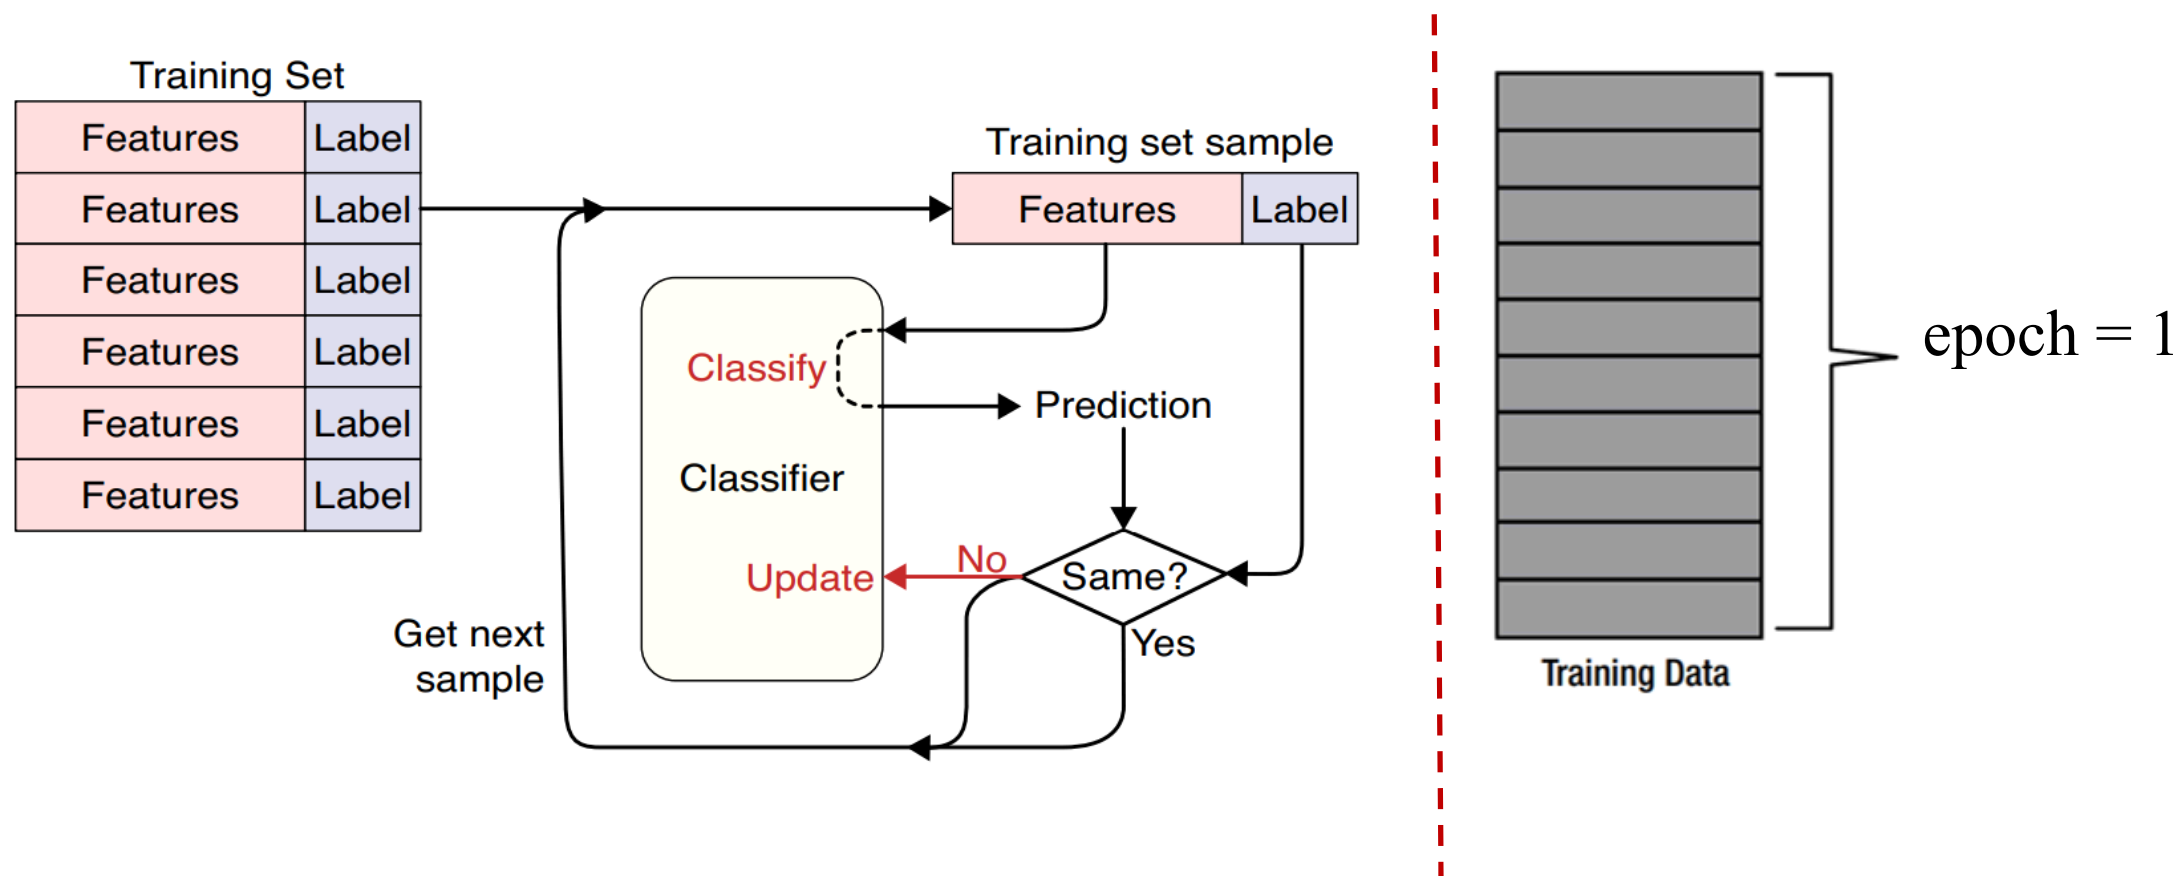
生醫光電所 吳育德

# Outline

- Leap from multiple linear regression to classification.
- Use a binary classifier to recognize a single digit in the MNIST dataset.
- Multiclass classification in recognizing all MNIST characters.

# The Essence of Training Classifier and Epoch

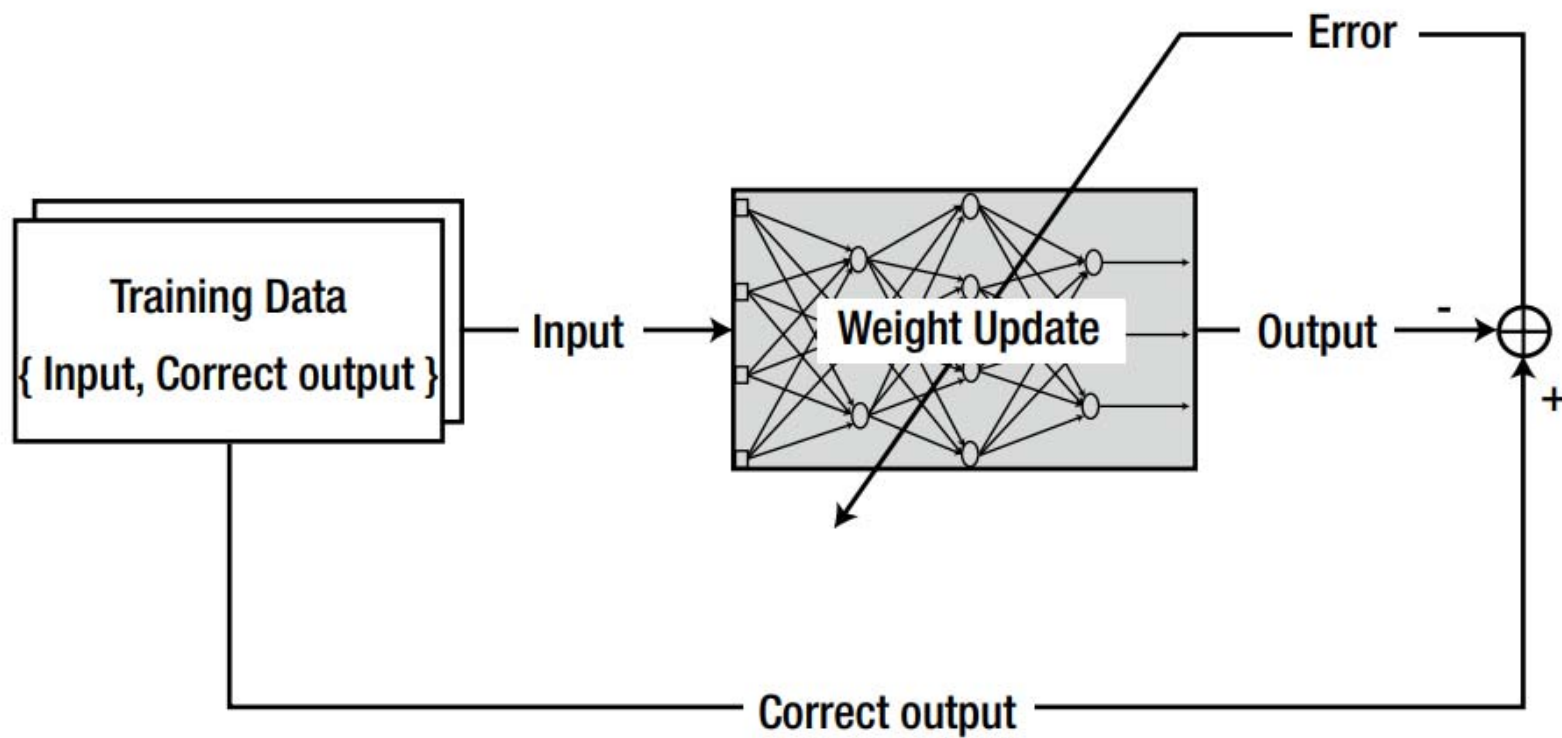
- Epoch is the **number of completed training cycles for all the training data.**



Andrew Glassner - Deep Learning\_ A Visual Approach-No Starch Press (2021)

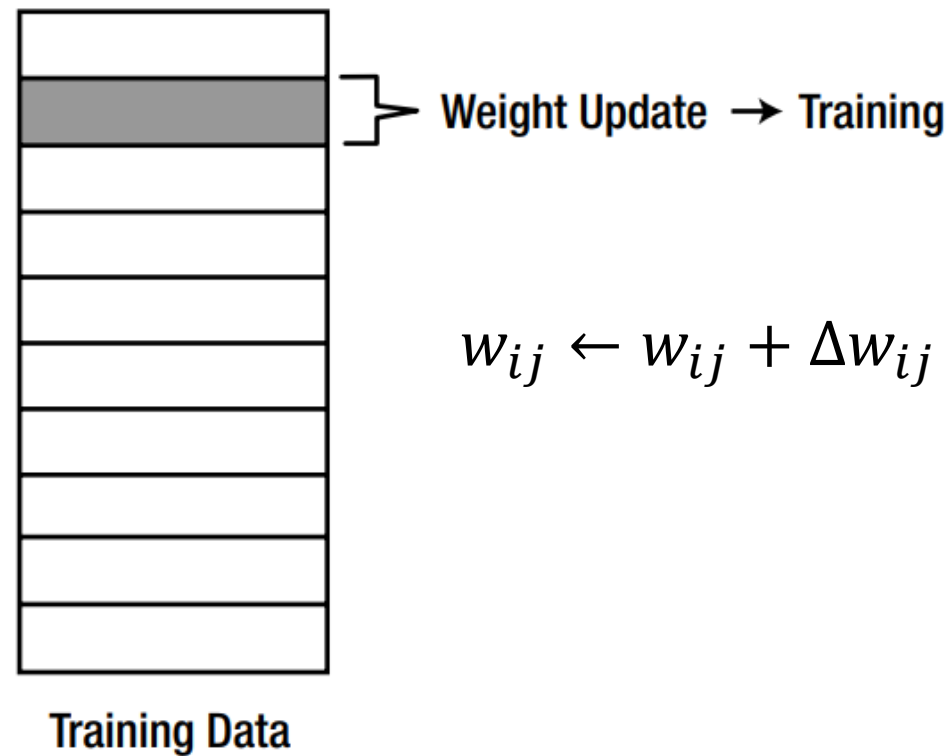
Phil Kim - MatLab Deep Learning with Machine Learning, Neural Networks and Artificial Intelligence 2017, Apress

# Supervised Learning of a Neural Network



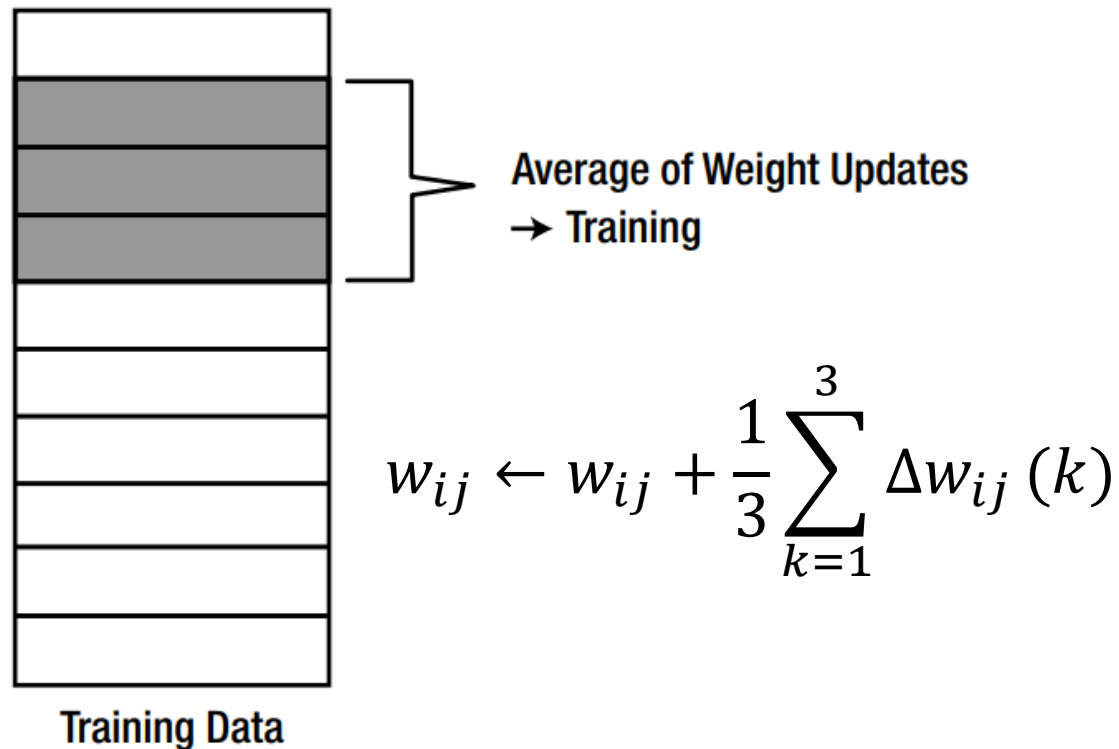
# Stochastic Gradient Descent (SGD) with Batch Size 1

- Assume 99 training data, the weights are updated **99 iterations** for each epoch
- Training **10 epochs** needs **990 iterations**




# Mini-batch Gradient descent with Batch Size 3

- Assume 99 training data, the weights are updated **33 iterations** for each epoch
- Training **10** epochs needs **330 iterations**



# Batch Gradient descent

- Assume 99 training data, the weights are updated 1 iteration for each epoch
- Training 10 epochs needs 10 iterations




# Where Linear Regression Fails: Binary Classification

- On busy nights, it's common for noisy customers to hang out in front of the pizzeria, until the neighbors eventually call the police.
- Roberto suspects that the same input variables such as temperature and tourists, also affect the number of loud customers by the entrance, and hence the likelihood of a police call.
- He wants to know in advance whether a police call is likely to happen, so that he can set up countermeasures—such as walk outside and beg people to lower their voices.

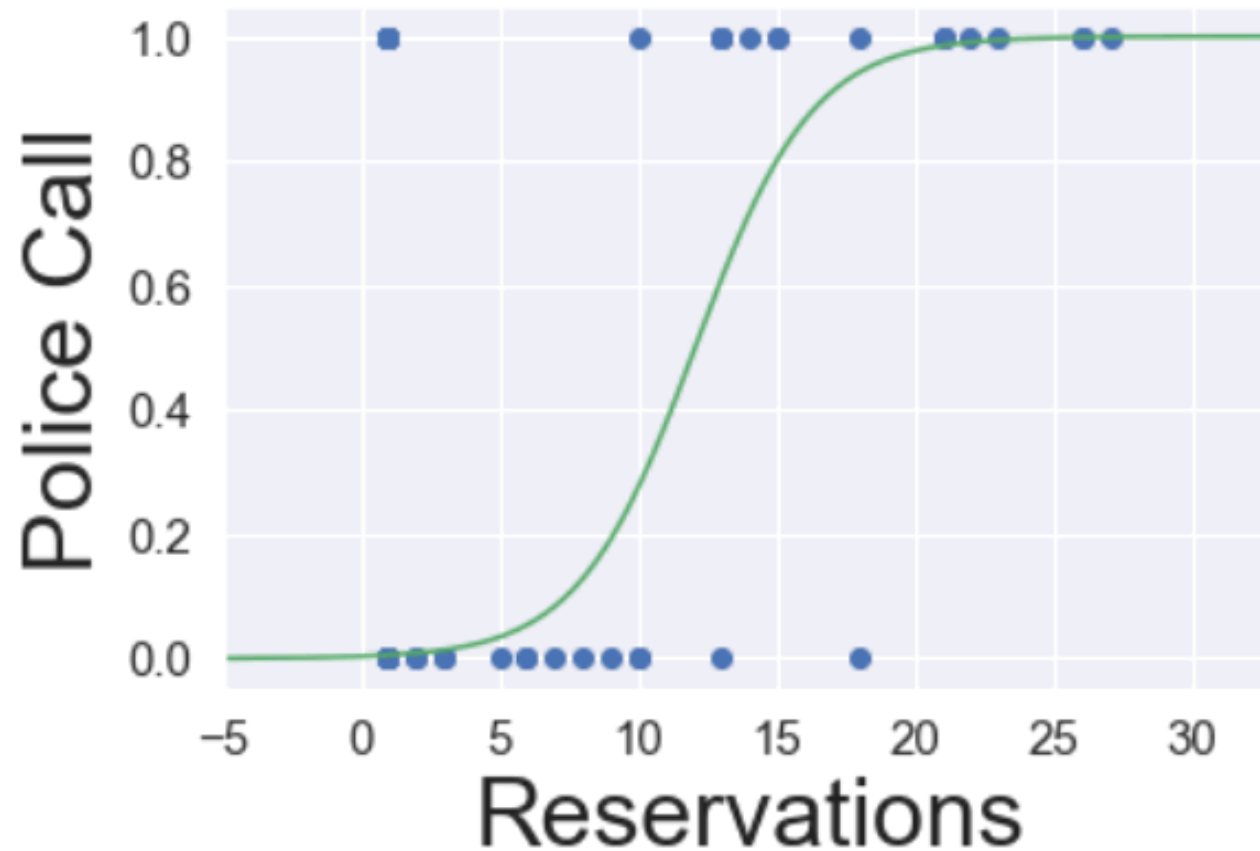
Reservations	Temperature	Tourists	Police
13	26	9	1
2	14	6	0
14	20	3	1
23	25	9	1
13	24	8	1
1	13	2	0
18	23	9	1
10	18	10	1
26	24	3	1
3	14	1	0
3	12	3	0
21	27	5	1
7	17	3	0
22	21	1	1
2	14	4	0
27	26	2	1
6	15	4	0
10	21	7	0
18	18	3	0
15	26	8	1
9	20	6	0
26	25	9	1
8	21	10	0
15	22	7	1
10	20	2	0
21	21	1	1
5	12	7	0
6	14	9	0
13	19	4	1
13	20	3	0




## “Reservations” column against the label



## “Reservations” column against the label




# Using the sigmoid function as an activation function



$$\frac{d(1 + e^{-x})^{-1}}{dx} = -(1 + e^{-x})^{-2}(-e^{-x}) = \frac{1}{1 + e^{-x}} \left( 1 - \frac{1}{1 + e^{-x}} \right)$$
$$\varphi'(x) = \varphi(x)(1 - \varphi(x))$$

# Multiple inputs and single output for Binary Classification



$$v = x_1 w_1 + x_2 w_2 + x_3 w_3 + b = \mathbf{x} \mathbf{w} + b$$

$$\text{where } \mathbf{x} = [x_1 \ x_2 \ x_3], \mathbf{w} = \begin{bmatrix} w_1 \\ w_2 \\ w_3 \end{bmatrix}$$

$y = \varphi(v)$ ,  $\varphi$  is the sigmoid activation function

**When  $b \equiv w_1 \neq 0$  and  $y = \varphi(v)$  with MSE**


$$\hat{v}_1 = [1 \ x_{12} \ x_{13} \ x_{14}] \begin{bmatrix} w_1 \\ w_2 \\ w_3 \\ w_4 \end{bmatrix} \Rightarrow e_1 = \varphi(\hat{v}_1) - y_1 = \hat{y}_1 - y_1$$

$$\vdots$$

$$\hat{v}_N = [1 \ x_{N2} \ x_{N3} \ x_{N4}] \begin{bmatrix} w_1 \\ w_2 \\ w_3 \\ w_4 \end{bmatrix} \Rightarrow e_N = \varphi(\hat{v}_N) - y_N = \hat{y}_N - y_N$$

$$\text{Loss}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N (\varphi(\hat{v}_n) - y_n)^2, \hat{\mathbf{y}}_{N \times 1} = \varphi([\mathbf{1}_{N \times 1} \ \mathbf{X}_{N \times 3}] \times \mathbf{w}_{4 \times 1})$$

# Mean Squared Error Has Too Many local Minimum



04\_plot\_losses.jpynb

# Minimize Mean Squared Error Using Gradient Descent

- The loss function for output

$$\text{Loss}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N (\varphi(\mathbf{w}_1 + x_{n2}\mathbf{w}_2 + x_{n3}\mathbf{w}_3 + x_{n4}\mathbf{w}_4) - y_n)^2 = \frac{1}{N} \sum_{n=1}^N e_N^2$$

- Minimize the loss function  $Loss$  w.r.t  $\mathbf{w}_1, \mathbf{w}_2, \mathbf{w}_3$  and  $\mathbf{w}_4$

$$\frac{\partial \text{Loss}}{\partial \mathbf{w}} = \begin{bmatrix} \frac{\partial \text{Loss}}{\partial \mathbf{w}_1} \\ \frac{\partial \text{Loss}}{\partial \mathbf{w}_2} \\ \frac{\partial \text{Loss}}{\partial \mathbf{w}_3} \\ \frac{\partial \text{Loss}}{\partial \mathbf{w}_4} \end{bmatrix} = \begin{bmatrix} \frac{\partial \text{Loss}}{\partial \hat{y}_n} \frac{\partial \hat{y}_n}{\partial \hat{v}_n} \frac{\partial \hat{v}_n}{\partial \mathbf{w}_1} \\ \frac{\partial \text{Loss}}{\partial \hat{y}_n} \frac{\partial \hat{y}_n}{\partial \hat{v}_n} \frac{\partial \hat{v}_n}{\partial \mathbf{w}_2} \\ \frac{\partial \text{Loss}}{\partial \hat{y}_n} \frac{\partial \hat{y}_n}{\partial \hat{v}_n} \frac{\partial \hat{v}_n}{\partial \mathbf{w}_3} \\ \frac{\partial \text{Loss}}{\partial \hat{y}_n} \frac{\partial \hat{y}_n}{\partial \hat{v}_n} \frac{\partial \hat{v}_n}{\partial \mathbf{w}_4} \end{bmatrix} \begin{bmatrix} \frac{2}{N} \sum_{n=1}^N e_n \varphi'(\hat{v}_n) \\ \frac{2}{N} \sum_{n=1}^N e_n \varphi'(\hat{v}_n) x_{n2} \\ \frac{2}{N} \sum_{n=1}^N e_n \varphi'(\hat{v}_n) x_{n3} \\ \frac{2}{N} \sum_{n=1}^N e_n \varphi'(\hat{v}_n) x_{n4} \end{bmatrix} = \frac{2}{N} [\mathbf{1}_{N \times 1} \quad \mathbf{X}_{N \times 3}]_{4 \times N}^T e_{N \times 1} \varphi'$$

- The **Batch Gradient Decent** method

$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} - \alpha \frac{\partial \text{Loss}}{\partial \mathbf{w}}$$

$$[\mathbf{1}_{N \times 1} \quad \mathbf{X}_{N \times 3}]_{4 \times N}^T \times e_{N \times 1} \varphi'$$

[illegible]

$x_2$

Reservations

13  
2  
14  
23  
13  
1  
18  
10  
26  
3  
3  
21  
7  
22  
2  
27  
6  
10  
18  
15  
9  
26  
8  
15  
10  
21  
5  
6  
13  
13

$x_3$

Temperature

26  
14  
20  
25  
24  
12  
23  
18  
24  
14  
12  
27  
17  
21  
12  
26  
15  
21  
18  
26  
20  
25  
21  
22  
20  
21  
12  
14  
19  
20

$x_4$  Pizzas 44 23 28 60 42 5 51 44 42 9 14 43 22 34 16 46 26 33 29 43 37 62 47 38 22 29 34 38 30 28

$$\begin{bmatrix} e_1 \varphi'(\hat{v}_1) \\ e_2 \varphi'(\hat{v}_2) \\ \vdots \\ \vdots \\ \vdots \\ e_{N-1} \varphi'(\hat{v}_{N-1}) \\ e_N \varphi'(\hat{v}_N) \end{bmatrix}$$



# Cross Entropy



- $L = -y \ln(\hat{y}) - (1 - y) \ln(1 - \hat{y})$  is the concatenation of the following two equations :

$$L = \begin{cases} -\ln(\hat{y}) & y = 1 \\ -\ln(1 - \hat{y}) & y = 0 \end{cases}$$

- Therefore, the cross entropy cost function often teams up with **sigmoid** and **softmax** activation functions in the neural network.

# Cross Entropy

- This cost function is proportional to the error.
- The cross entropy function is much more sensitive to the error than quadratic function.



$$\begin{aligned}
\frac{\partial L}{\partial \hat{y}} &= \frac{\partial(-y \ln(\hat{y}) - (1-y) \ln(1-\hat{y}))}{\partial \hat{y}} \\
&= -y \frac{1}{\hat{y}} - (1-y) \frac{-1}{1-\hat{y}} \\
&= \frac{-y(1-\hat{y}) + (1-y)\hat{y}}{\hat{y}(1-\hat{y})} \\
&= \frac{-(y-\hat{y})}{\hat{y}(1-\hat{y})}
\end{aligned}$$

$$\hat{y} = \varphi(\hat{v})$$

$$\frac{\partial \hat{y}}{\partial \hat{v}} = \varphi(\hat{v})(1 - \varphi(\hat{v})) = \hat{y}(1 - \hat{y})$$

$$\frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial \hat{v}} = \frac{-(y-\hat{y})}{\hat{y}(1-\hat{y})} \times \hat{y}(1-\hat{y}) = -(y-\hat{y}) = \hat{y} - y = \varphi(\hat{v}) - y \rightarrow e$$

**When  $b \equiv w_1 \neq 0$  and  $y = \varphi(v)$  with CE loss**

$$\hat{v}_1 = [1 \ x_{\textcolor{red}{1}2} \ x_{\textcolor{red}{1}3} \ x_{\textcolor{red}{1}4}] \begin{bmatrix} w_1 \\ w_2 \\ w_3 \\ w_4 \end{bmatrix} \Rightarrow \textcolor{blue}{e}_1 = \varphi(\hat{v}_1) - \textcolor{blue}{y}_1 = \hat{y}_1 - \textcolor{blue}{y}_1$$


$$\vdots$$

$$\hat{v}_N = [1 \ x_{\textcolor{red}{N}2} \ x_{\textcolor{red}{N}3} \ x_{\textcolor{red}{N}4}] \begin{bmatrix} w_1 \\ w_2 \\ w_3 \\ w_4 \end{bmatrix} \Rightarrow \textcolor{blue}{e}_N = \varphi(\hat{v}_N) - \textcolor{blue}{y}_N = \hat{y}_N - \textcolor{blue}{y}_N$$

$$\text{Loss}(\mathbf{w}) = \frac{-1}{N} \sum_{n=1}^N (y_n \ln(\hat{y}_n) + (1 - y_n) \ln(1 - \hat{y}_n)),$$

$$\hat{\mathbf{y}}_{N \times 1} = \varphi([\mathbf{1}_{N \times 1} \ \mathbf{X}_{N \times 3}] \times \mathbf{w}_{4 \times 1})$$

# Cross Entropy Error is Nice and Smooth



04\_plot\_losses.jpynb

# Minimize Mean Cross Entropy Using Gradient Descent

- The loss function for output

$$\text{Loss}(\mathbf{w}) = \frac{-1}{N} \sum_{n=1}^N (y_n \ln(\hat{y}_n) + (1 - y_n) \ln(1 - \hat{y}_n))$$

- Minimize the loss function  $\text{Loss}$  w.r.t  $\mathbf{w}_1, \mathbf{w}_2, \mathbf{w}_3$  and  $\mathbf{w}_4$

$$\frac{\partial \text{Loss}}{\partial \mathbf{w}} = \begin{bmatrix} \frac{\partial \text{Loss}}{\partial \mathbf{w}_1} \\ \frac{\partial \text{Loss}}{\partial \mathbf{w}_2} \\ \frac{\partial \text{Loss}}{\partial \mathbf{w}_3} \\ \frac{\partial \text{Loss}}{\partial \mathbf{w}_4} \end{bmatrix} = \begin{bmatrix} \frac{\partial \text{Loss}}{\partial \hat{y}_n} \frac{\partial \hat{y}_n}{\partial \hat{v}_n} \frac{\partial \hat{v}_n}{\partial \mathbf{w}_1} \\ \frac{\partial \text{Loss}}{\partial \hat{y}_n} \frac{\partial \hat{y}_n}{\partial \hat{v}_n} \frac{\partial \hat{v}_n}{\partial \mathbf{w}_2} \\ \frac{\partial \text{Loss}}{\partial \hat{y}_n} \frac{\partial \hat{y}_n}{\partial \hat{v}_n} \frac{\partial \hat{v}_n}{\partial \mathbf{w}_3} \\ \frac{\partial \text{Loss}}{\partial \hat{y}_n} \frac{\partial \hat{y}_n}{\partial \hat{v}_n} \frac{\partial \hat{v}_n}{\partial \mathbf{w}_4} \end{bmatrix} = \begin{bmatrix} \frac{1}{N} \sum_{n=1}^N e_n \\ \frac{1}{N} \sum_{n=1}^N e_n x_{n2} \\ \frac{1}{N} \sum_{n=1}^N e_n x_{n3} \\ \frac{1}{N} \sum_{n=1}^N e_n x_{n4} \end{bmatrix} = \frac{2}{N} [\mathbf{1}_{N \times 1} \quad \mathbf{X}_{N \times 3}]_{4 \times N}^T e_{N \times 1}$$

- The **Batch Gradient Decent** method

$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} - \alpha \frac{\partial \text{Loss}}{\partial \mathbf{w}}$$

$$[\mathbf{1}_{N \times 1} \quad \mathbf{X}_{N \times 3}]_{4 \times N}^T \times e_{N \times 1}$$

[illegible]

$x_2$

Reservations
13
2
14
23
13
1
18
10
26
3
3
21
7
22
2
27
6
10
18
15
9
26
8
15
10
21
5
6
13
13

$x_3$  Temperature

$x_4$  Pizzas 44 23 28 60 42 5 51 44 42 9 14 43 22 34 16 46 26 33 29 43 37 62 47 38 22 29 34 38 30 28

$$\begin{bmatrix} e_1 = \varphi(\hat{v}_1) - y_1 \\ e_2 = \varphi(\hat{v}_2) - y_2 \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ e_N = \varphi(\hat{v}_N) - y_N \end{bmatrix}$$

```
import numpy as np
```

```
def sigmoid(z):  
    return 1 / (1 + np.exp(-z))
```

```
def forward(X, w):  
    weighted_sum = np.matmul(X, w)  
    return sigmoid(weighted_sum)
```

```
def classify(X, w):  
    return np.round(forward(X, w))
```

```
def loss(X, Y, w):  
    y_hat = forward(X, w)  
    first_term = Y * np.log(y_hat)  
    second_term = (1 - Y) * np.log(1 - y_hat)  
    return -np.average(first_term + second_term)
```

```
def gradient(X, Y, w):  
    return np.matmul(X.T, (forward(X, w) - Y)) / X.shape[0]
```



```
def train(X, Y, iterations, lr):  
    w = np.zeros((X.shape[1], 1))  
    for i in range(iterations):  
        print("Iteration %4d => Loss: %.20f" % (i, loss(X, Y, w)))  
        w -= gradient(X, Y, w) * lr  
    return w
```

```
def test(X, Y, w):  
    total_examples = X.shape[0]  
    correct_results = np.sum(classify(X, w) == Y)  
    success_percent = correct_results * 100 / total_examples  
    print("\nSuccess: %d/%d (%.2f%%)" %  
          (correct_results, total_examples, success_percent))
```

*# Prepare data*

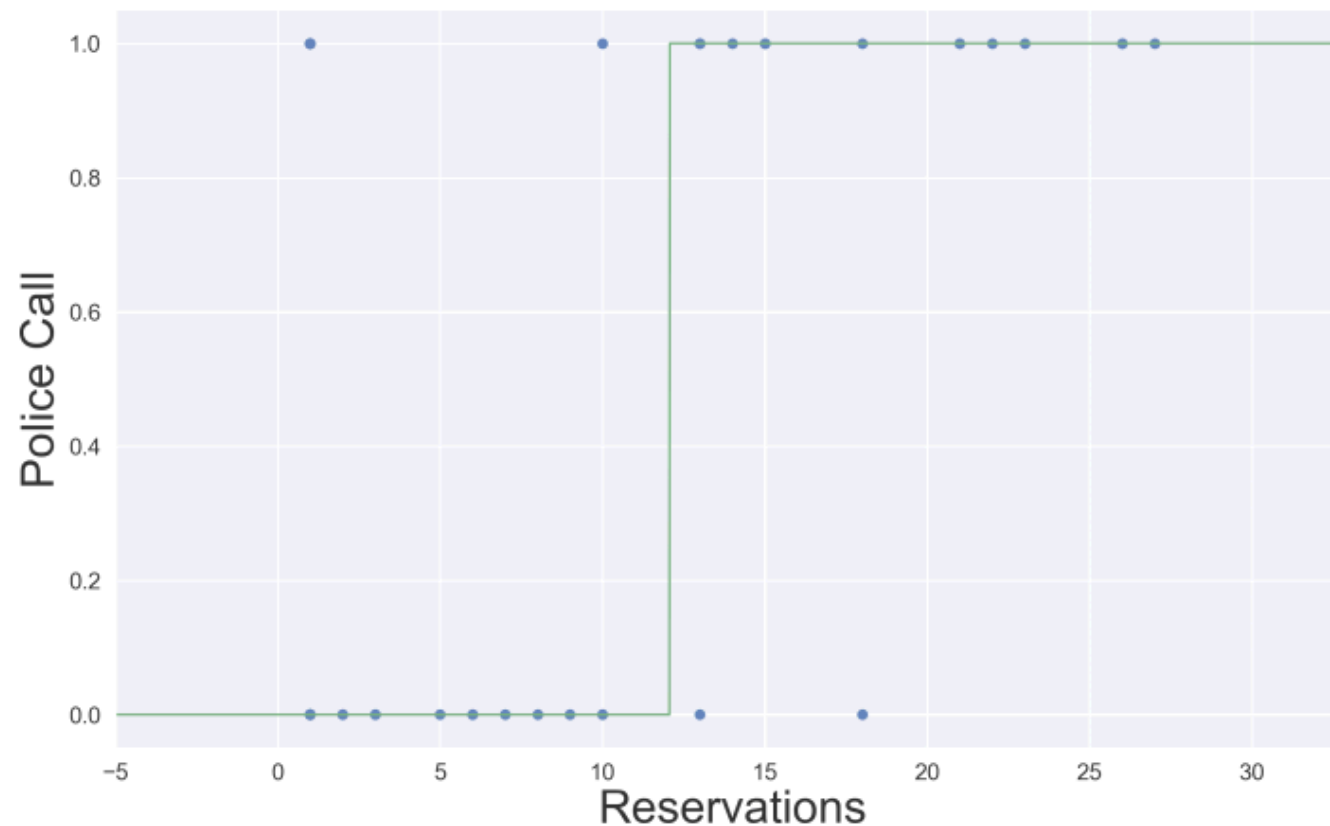
```
x1, x2, x3, y = np.loadtxt("police.txt", skiprows=1, unpack=True)  
X = np.column_stack((np.ones(x1.size), x1, x2, x3))  
Y = y.reshape(-1, 1)  
w = train(X, Y, iterations=10000, lr=0.001)
```

*# Test it*

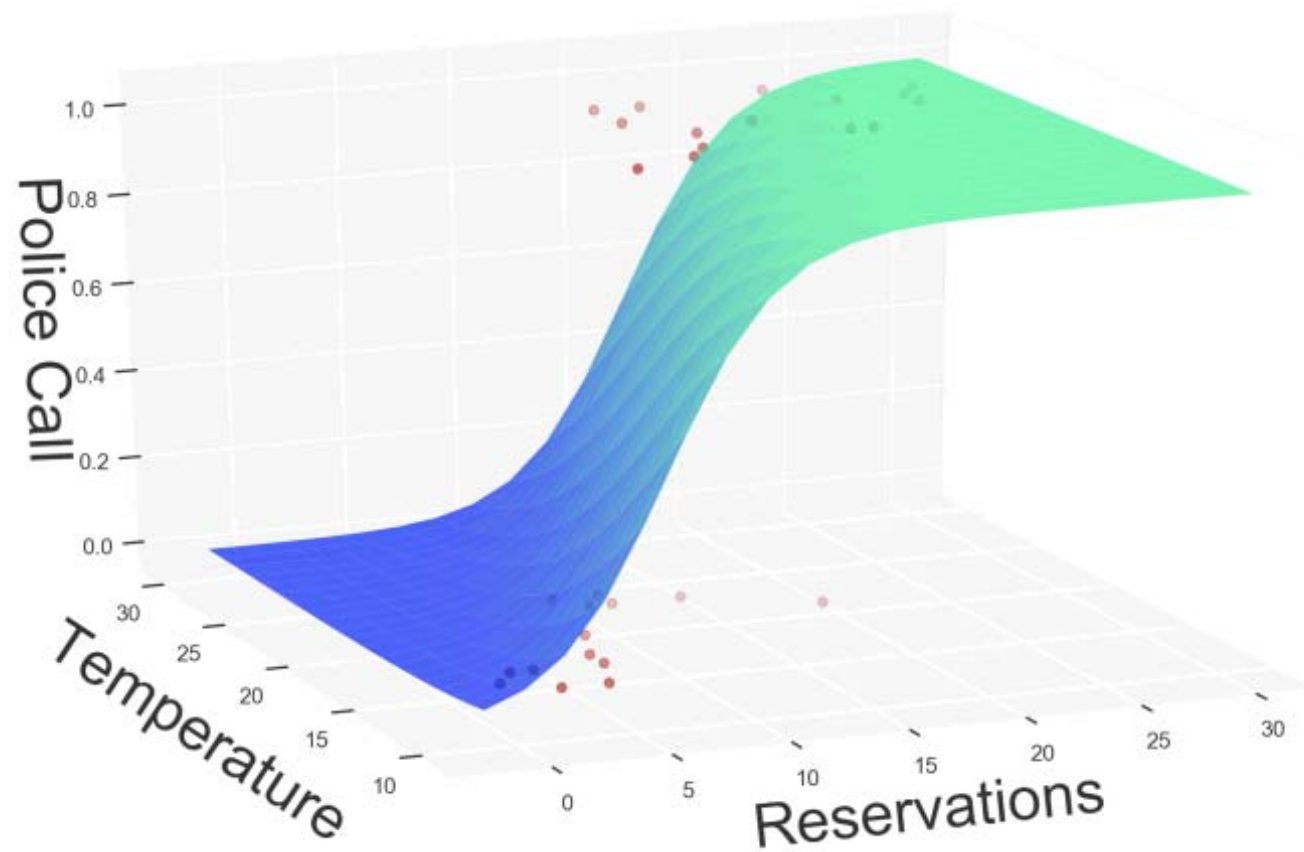
```
test(X, Y, w)
```

# Binary Classification

Below a certain threshold around 12, the value of the model is under 0.5, and `classify()` rounds it to 0. Above that threshold, the value of the model is over 0.5, and `classify()` rounds it to 1. The result is a stepwise model function.




# Visualization of Classifier's Model



# Binary Classification of MNIST




- Digits are made up of 28 by 28 grayscale pixels, each represented by one byte.
- In MNIST's grayscale, 0 stands for “perfect background white,” and 255 stands for “perfect foreground black.”
- It contains 70,000 examples, neatly partitioned into 7,000 examples for each digit from 0 to 9.





# MNIST Training and Testing Data

X



Y



---

```
import numpy as np
import gzip
import struct

def load_images(filename):
    # Open and unzip the file of images:
    with gzip.open(filename, 'rb') as f:
        # Read the header information into a bunch of variables
        _ignored, n_images, columns, rows = struct.unpack('>IIII', f.read(16))
        # Read all the pixels into a NumPy array of bytes:
        all_pixels = np.frombuffer(f.read(), dtype=np.uint8)
        # Reshape the pixels into a matrix where each line is an image:
        return all_pixels.reshape(n_images, columns * rows)

def prepend_bias(X):
    # Insert a column of 1s in the position 0 of X.
    # ("axis=1" stands for: "insert a column, not a row")
    return np.insert(X, 0, 1, axis=1)

# 60000 images, each 785 elements (1 bias + 28 * 28 pixels)
X_train = prepend_bias(load_images("../data/mnist/train-images-idx3-ubyte.gz"))

# 10000 images, each 785 elements, with the same structure as X_train
X_test = prepend_bias(load_images("../data/mnist/t10k-images-idx3-ubyte.gz"))
```



```

def load_labels(filename):
    # Open and unzip the file of images:
    with gzip.open(filename, 'rb') as f:
        # Skip the header bytes:
        f.read(8)
        # Read all the labels into a list:
        all_labels = f.read()
        # Reshape the list of labels into a one-column matrix:
        return np.frombuffer(all_labels, dtype=np.uint8).reshape(-1, 1)

def encode_fives(Y):
    # Convert all 5s to 1, and everything else to 0
    return (Y == 5).astype(int)


# 60K labels, each with value 1 if the digit is a five, and 0 otherwise
Y_train = encode_fives(load_labels("../data/mnist/train-labels-idx1-ubyte.gz"))

# 10000 labels, with the same encoding as Y_train
Y_test = encode_fives(load_labels("../data/mnist/t10k-labels-idx1-ubyte.gz"))

```

reshape(-1, 1) means: “Arrange these data into a matrix with one column, and however many rows you need.”

# Multiple inputs and single output for Binary Classification



$$v = x_1 w_1 + x_2 w_2 + x_3 w_3 + \cdots x_{785} w_{785} = \mathbf{x} \mathbf{w}$$

$$\text{where } \mathbf{x} = [x_1 \ \cdots \ x_{785}], \mathbf{w} = \begin{bmatrix} w_1 \\ \vdots \\ w_{785} \end{bmatrix}$$

$y = \varphi(v)$ ,  $\varphi$  is the sigmoid activation function



**When  $b \equiv w_1 \neq 0$  and  $\mathbf{y} = \varphi(\mathbf{v})$  with CE loss**

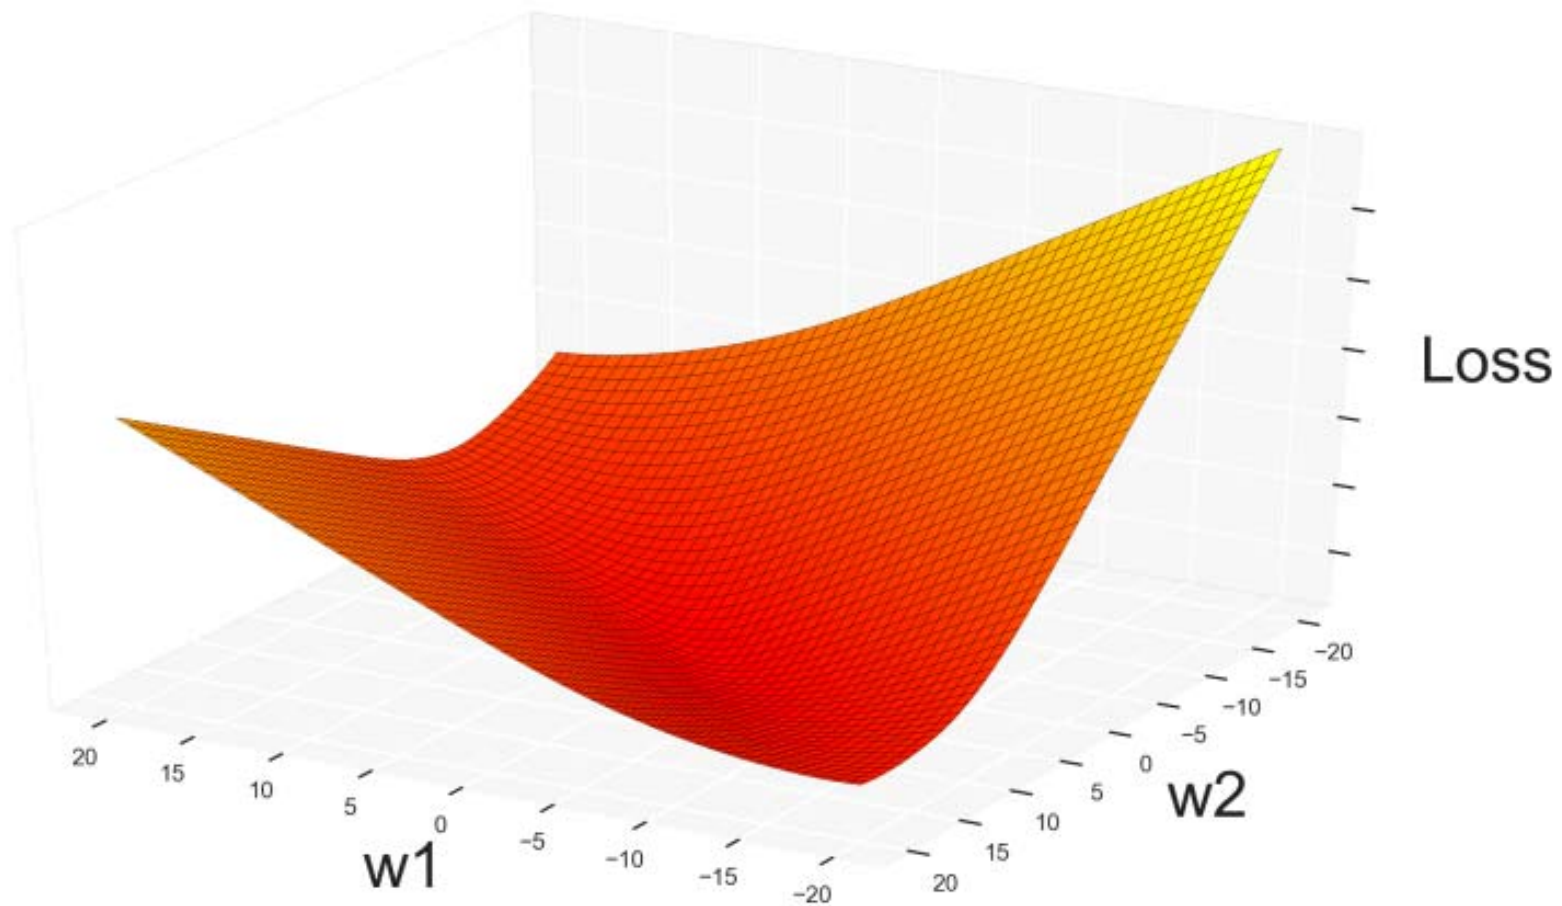
$$\hat{v}_1 = [1 \ x_{\mathbf{1}2} \ \cdots \ x_{\mathbf{1}785}] \begin{bmatrix} w_1 \\ w_2 \\ \vdots \\ w_{785} \end{bmatrix} \Rightarrow e_1 = \varphi(\hat{v}_1) - y_1 = \hat{y}_1 - y_1$$

$$\hat{v}_N = [1 \ x_{\mathbf{N}2} \ \cdots \ x_{\mathbf{N}785}] \begin{bmatrix} w_1 \\ w_2 \\ \vdots \\ w_{785} \end{bmatrix} \Rightarrow e_N = \varphi(\hat{v}_N) - y_N = \hat{y}_N - y_N$$

$$\text{Loss}(\mathbf{w}) = \frac{-1}{N} \sum_{n=1}^N (y_n \ln(\hat{y}_n) + (1 - y_n) \ln(1 - \hat{y}_n)),$$

$$\hat{\mathbf{y}}_{N \times 1} = \varphi([\mathbf{1}_{N \times 1} \ \mathbf{X}_{N \times 784}] \times \mathbf{w}_{785 \times 1})$$

# Cross Entropy Error is Nice and Smooth



04\_plot\_losses.jpynb

# Minimize Mean Cross Entropy Using Gradient Descent

- The loss function for output

$$\text{Loss}(\mathbf{w}) = \frac{-1}{N} \sum_{n=1}^N (y_n \ln(\hat{y}_n) + (1 - y_n) \ln(1 - \hat{y}_n))$$

- Minimize the loss function  $\text{Loss}$  w.r.t  $\mathbf{w}_1, \mathbf{w}_2, \dots, \mathbf{w}_{785}$

$$\frac{\partial \text{Loss}}{\partial \mathbf{w}} = \begin{bmatrix} \frac{\partial \text{Loss}}{\partial \mathbf{w}_1} \\ \frac{\partial \text{Loss}}{\partial \mathbf{w}_2} \\ \vdots \\ \frac{\partial \text{Loss}}{\partial \mathbf{w}_{785}} \end{bmatrix} = \begin{bmatrix} \frac{1}{N} \sum_{n=1}^N e_n \\ \frac{1}{N} \sum_{n=1}^N e_n x_{n2} \\ \vdots \\ \frac{1}{N} \sum_{n=1}^N e_n x_{n785} \end{bmatrix} = \frac{1}{N} [\mathbf{1}_{N \times 1} \quad \mathbf{X}_{N \times 784}]_{785 \times N}^T e_{N \times 1}$$

- The **Batch Gradient Decent** method

$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} - \alpha \frac{\partial \text{Loss}}{\partial \mathbf{w}}$$

$$[\mathbf{1}_{N \times 1} \quad \mathbf{X}_{N \times 3}]_{785 \times N}^T \times e_{N \times 1}$$


$x_1$

1的列向量 共有 N 個

$x_2$

N 個影像中左上角第一個畫素灰階形成的列向量

$\vdots$



$x_{785}$

N 個影像中右下角第一個畫素灰階形成的列向量

$$\begin{bmatrix} e_1 = \varphi(\hat{v}_1) - y_1 \\ e_2 = \varphi(\hat{v}_2) - y_2 \\ \vdots \\ e_N = \varphi(\hat{v}_N) - y_N \end{bmatrix}$$

```
import numpy as np
```

```
def sigmoid(z):  
    return 1 / (1 + np.exp(-z))
```

```
def forward(X, w):  
    weighted_sum = np.matmul(X, w)  
    return sigmoid(weighted_sum)
```

```
def classify(X, w):  
    return np.round(forward(X, w))
```

```
def loss(X, Y, w):  
    y_hat = forward(X, w)  
    first_term = Y * np.log(y_hat)  
    second_term = (1 - Y) * np.log(1 - y_hat)  
    return -np.average(first_term + second_term)
```

```
def gradient(X, Y, w):  
    return np.matmul(X.T, (forward(X, w) - Y)) / X.shape[0]
```

```


def train(X, Y, iterations, lr):
    w = np.zeros((X.shape[1], 1))
    for i in range(iterations):
        print("Iteration %4d => Loss: %.20f" % (i, loss(X, Y, w)))
        w -= gradient(X, Y, w) * lr
    return w

def test(X, Y, w):
    total_examples = X.shape[0]
    correct_results = np.sum(classify(X, w) == Y)
    success_percent = correct_results * 100 / total_examples
    print("\nSuccess: %d/%d (%.2f%%)" %
          (correct_results, total_examples, success_percent))

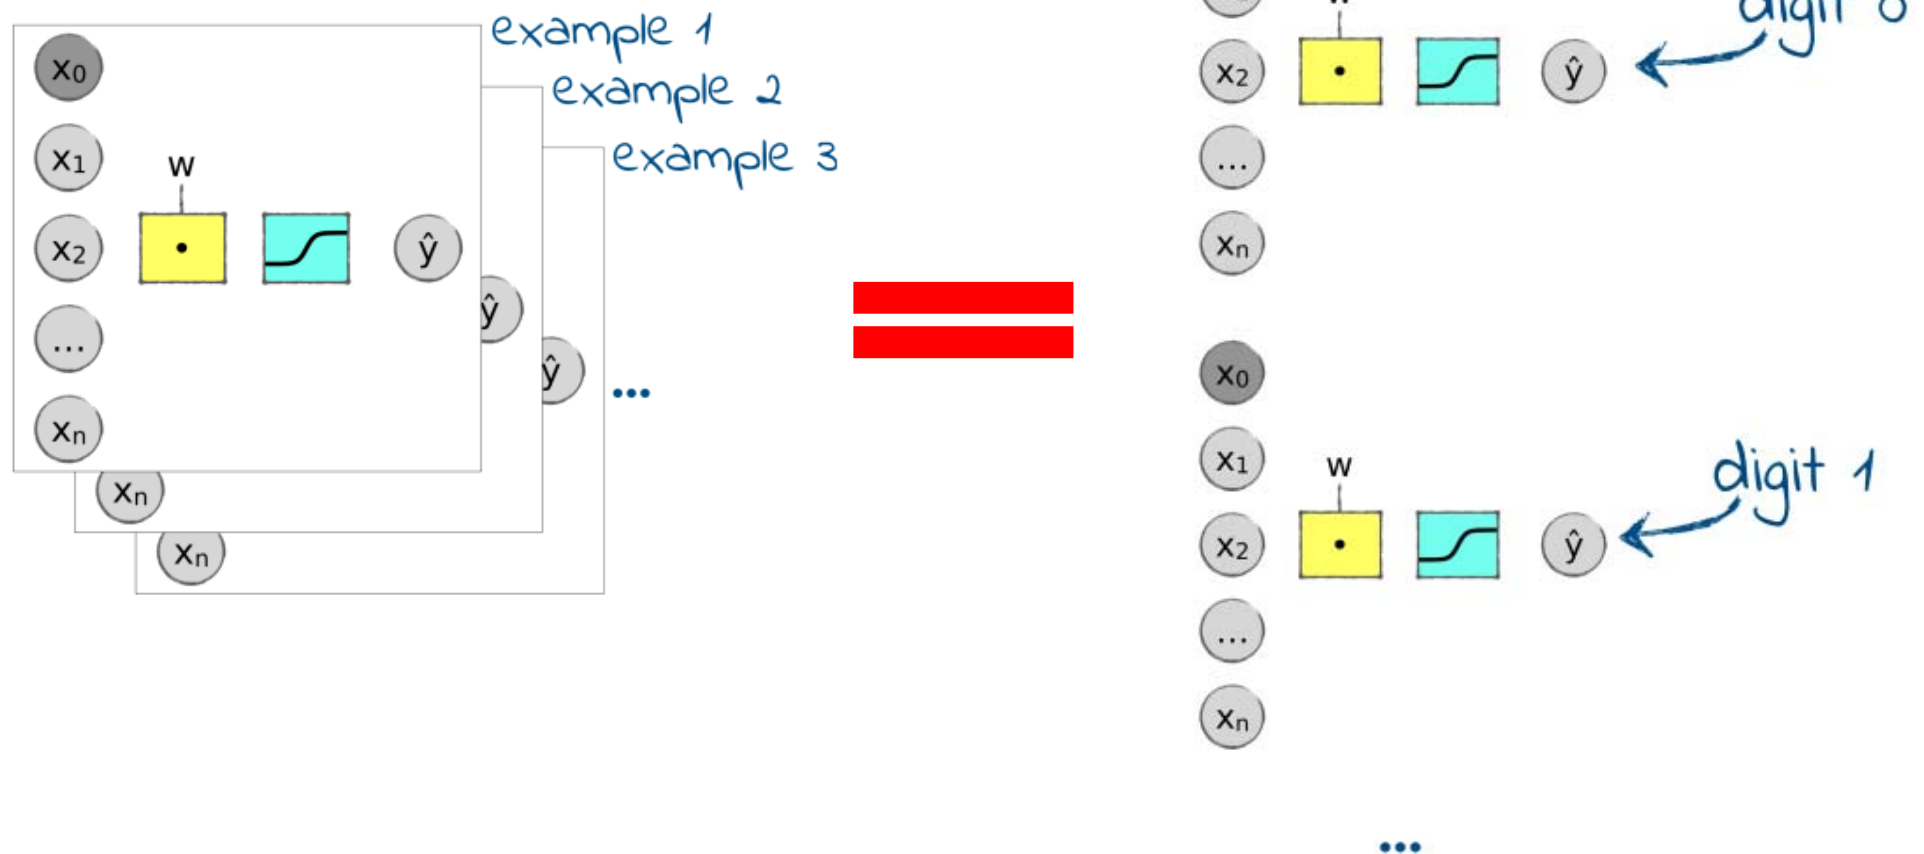
w = train(data.X_train, data.Y_train, iterations=100, lr=1e-5)
test(data.X_test, data.Y_test, w)

```

# Multiple Classification Using Sigmoid




# Multiple Classification Using Sigmoid





# Multiple inputs and outputs for Multiple Classification



$$\hat{v}_{i,j} = x_{i,1}w_{1,j} + x_{i,2}w_{2,j} + x_{i,3}w_{3,j} + \cdots x_{i,785}w_{785,j} = \mathbf{x}_i \mathbf{w}_j,$$

$$\text{where } \mathbf{x}_i = [x_{i,1} \cdots x_{i,785}], \mathbf{w}_j = \begin{bmatrix} w_{1,j} \\ \vdots \\ w_{785,j} \end{bmatrix}, i = 1, \cdots, N, j = 1, \cdots, 10$$

$$\hat{y}_{i,j} = \varphi(v_{i,j}), \varphi \text{ is the sigmoid activation function}$$

# One-Hot Encoding

A possible output  $\hat{y}_1, \dots, \hat{y}_{10}$

"0"	"1"	"2"	"3"	"4"	"5"	"6"	"7"	"8"	"9"
0.111	0.005	0.787	0.170	0.001	0.176	0.352	0.001	0.073	0.003

## Labels

4

The diagram shows a vertical vector on the left with elements 3, 5, 3, 0, and an ellipsis. An arrow points to a 5x10 grid on the right. The grid contains 0s and 1s. The first row has a 1 in the 4th column. The second row has a 1 in the 6th column. The third row has a 1 in the 4th column. The fourth row has a 1 in the 1st column. The fifth row has an ellipsis in every column. The result of the dot product is 18.

---

```
def one_hot_encode(Y):  
    n_labels = Y.shape[0]  
    n_classes = 10  
    encoded_Y = np.zeros((n_labels, n_classes))  
    for i in range(n_labels):  
        label = Y[i]  
        encoded_Y[i][label] = 1  
    return encoded_Y
```

- `one_hot_encode()` initializes a matrix of zeros with one row per label, and one column per class
- `Y.shape[0]` means “the number of rows in Y”
- Then it walks through the matrix, flipping the “hot” values to 1.

```
# 60K labels, each a single digit from 0 to 9
```

```
Y_train_unencoded = load_labels("../data/mnist/train-labels-idx1-ubyte.gz")
```

```
# 60K labels, each consisting of 10 one-hot encoded elements
```

```
Y_train = one_hot_encode(Y_train_unencoded)
```

```
# 10000 labels, each a single digit from 0 to 9
```

```
Y_test = load_labels("../data/mnist/t10k-labels-idx1-ubyte.gz")
```

**When  $b \equiv w_1 \neq 0$  and  $\hat{y} = \varphi(\hat{v})$  with CE loss**

$$\begin{bmatrix} 1 & x_{\mathbf{1},2} & \cdots & x_{\mathbf{1},785} \\ 1 & x_{\mathbf{2},2} & \cdots & x_{\mathbf{2},785} \\ \vdots & \vdots & \vdots & \vdots \\ 1 & x_{\mathbf{N},2} & \cdots & x_{\mathbf{N},785} \end{bmatrix} \begin{bmatrix} w_{1,1} & w_{1,2} & \cdots & w_{1,10} \\ w_{2,1} & w_{2,2} & \cdots & w_{2,10} \\ \vdots & \vdots & \vdots & \vdots \\ w_{785,1} & w_{785,2} & \cdots & w_{785,10} \end{bmatrix} = \begin{bmatrix} \hat{v}_{1,1} & \cdots & \hat{v}_{1,10} \\ \hat{v}_{2,1} & \cdots & \hat{v}_{2,10} \\ \vdots & \vdots & \vdots \\ \hat{v}_{N,1} & \cdots & \hat{v}_{N,10} \end{bmatrix}_{N \times 10}$$

$$\Rightarrow \hat{\mathbf{v}}_{N \times 10} = [\mathbf{1}_{N \times 1} \quad \mathbf{X}_{N \times 784}] \times \mathbf{w}_{785 \times 10}$$

$$\Rightarrow \hat{\mathbf{y}}_{N \times 10} = \varphi([\mathbf{1}_{N \times 1} \quad \mathbf{X}_{N \times 784}] \times \mathbf{w}_{785 \times 10})$$

$$\Rightarrow e_{1,j} = \varphi(\hat{v}_{1,j}) - y_{1,j} = \hat{y}_{1,j} - y_{1,j}, j = 1, \cdots 10$$

$$\vdots$$

$$\Rightarrow e_{N,j} = \varphi(\hat{v}_{N,j}) - y_{N,j} = \hat{y}_{N,j} - y_{N,j}, j = 1, \cdots 10$$

$$\Rightarrow \hat{\mathbf{e}}_{N \times 10} = \varphi([\mathbf{1}_{N \times 1} \quad \mathbf{X}_{N \times 784}] \times \mathbf{w}_{785 \times 10}) - \mathbf{Y}_{N \times 10}$$

# Minimize Mean Cross Entropy Using Gradient Descent

- The loss function for output

$$\text{Loss}(\mathbf{w}) = \sum_{j=1}^{10} \left\{ \frac{-1}{N} \sum_{n=1}^N (y_{n,j} \ln(\hat{y}_{n,j}) + (1 - y_{n,j}) \ln(1 - \hat{y}_{n,j})) \right\}$$

Minimize  $\text{Loss}(\mathbf{w})$  w.r.t  $w_{1,1}, w_{2,1}, \dots, w_{785,1}, \dots, w_{1,10}, w_{2,10}, \dots, w_{785,10}$

$$\begin{aligned} \frac{\partial \text{Loss}}{\partial \mathbf{w}} &= \begin{bmatrix} \frac{\partial \text{Loss}}{\partial w_{1,1}} & \frac{\partial \text{Loss}}{\partial w_{1,2}} & \dots & \frac{\partial \text{Loss}}{\partial w_{1,10}} \\ \frac{\partial \text{Loss}}{\partial w_{2,1}} & \frac{\partial \text{Loss}}{\partial w_{2,2}} & \dots & \frac{\partial \text{Loss}}{\partial w_{2,10}} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial \text{Loss}}{\partial w_{785,1}} & \frac{\partial \text{Loss}}{\partial w_{785,2}} & \dots & \frac{\partial \text{Loss}}{\partial w_{785,10}} \end{bmatrix} = \begin{bmatrix} \frac{1}{N} \sum_{n=1}^N e_{n,1} & \frac{1}{N} \sum_{n=1}^N e_{n,2} & \dots & \frac{1}{N} \sum_{n=1}^N e_{n,10} \\ \frac{1}{N} \sum_{n=1}^N e_{n,1} x_{n,2} & \frac{1}{N} \sum_{n=1}^N e_{n,2} x_{n,2} & \dots & \frac{1}{N} \sum_{n=1}^N e_{n,10} x_{n,2} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{1}{N} \sum_{n=1}^N e_{n,1} x_{n,785} & \frac{1}{N} \sum_{n=1}^N e_{n,2} x_{n,785} & \dots & \frac{1}{N} \sum_{n=1}^N e_{n,10} x_{n,785} \end{bmatrix} \\ &= \frac{1}{N} [\mathbf{1}_{N \times 1} \quad \mathbf{X}_{N \times 284}]_{785 \times N}^T \mathbf{e}_{N \times 10} \end{aligned}$$

- The steepest (gradient) decent method

$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} - \alpha \frac{\partial \text{Loss}}{\partial \mathbf{w}}$$

---

```
import numpy as np
```

```
def sigmoid(z):  
    return 1 / (1 + np.exp(-z))
```

```
def forward(X, w):  
    weighted_sum = np.matmul(X, w)  
    return sigmoid(weighted_sum)
```

`weighted_sum.shape` → (N, 10)

```
def classify(X, w):  
    y_hat = forward(X, w)  
    labels = np.argmax(y_hat, axis=1)  
    return labels.reshape(-1, 1)
```

`labels.shape` → (N, 1)

```
def loss(X, Y, w):  
    y_hat = forward(X, w)  
    first_term = Y * np.log(y_hat)  
    second_term = (1 - Y) * np.log(1 - y_hat)  
    return -np.sum(first_term + second_term) / X.shape[0]
```

→ (N, 10)

→ (N, 10)

```
def gradient(X, Y, w):  
    return np.matmul(X.T, (forward(X, w) - Y)) / X.shape[0]
```


→ (785, 10)

```
def report(iteration, X_train, Y_train, X_test, Y_test, w):  
    matches = np.count_nonzero(classify(X_test, w) == Y_test)  
    n_test_examples = Y_test.shape[0]  
    matches = matches * 100.0 / n_test_examples  
    training_loss = loss(X_train, Y_train, w)  
    print("%d - Loss: %.20f, %.2f%%" % (iteration, training_loss, matches))
```

```
def train(X_train, Y_train, X_test, Y_test, iterations, lr):  
    w = np.zeros((X_train.shape[1], Y_train.shape[1]))  
    for i in range(iterations):  
        report(i, X_train, Y_train, X_test, Y_test, w)  
        w -= gradient(X_train, Y_train, w) * lr  
    report(iterations, X_train, Y_train, X_test, Y_test, w)  
    return w
```

```
w = train(X_train, Y_train,  
          X_test, Y_test,  
          iterations=200, lr=1e-5)
```

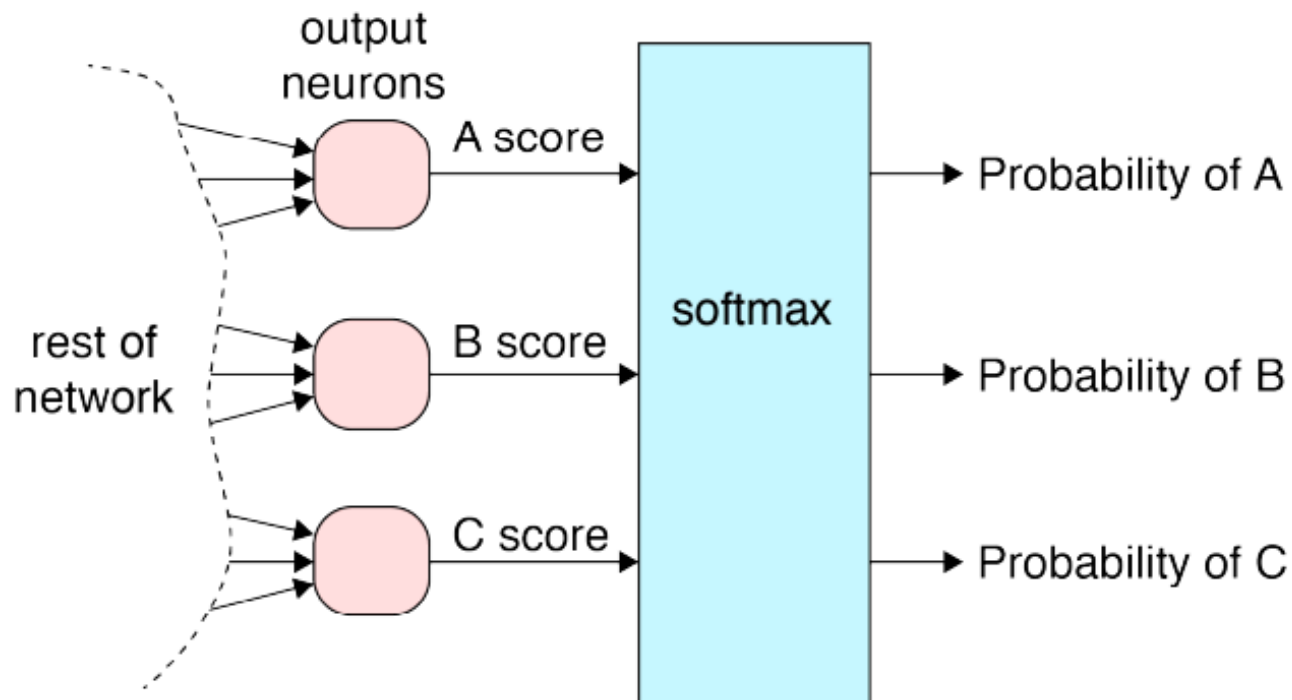
# Multiple Classification Using **Softmax**





# The softmax function turns all the scores into probabilities

In general, multiclass classifiers employ the softmax function as the activation function of the output node.



# The softmax function turns all the scores into probabilities

- The output from the  $i$ -th output node of the softmax function is :

$$y_i = \sigma(v_i) = \frac{e^{v_i}}{e^{v_1} + e^{v_2} + e^{v_3} + \dots + e^{v_M}} = \frac{e^{v_i}}{\sum_{k=1}^M e^{v_k}}$$

$v_i$  is the weighted sum of the  $i$ -th output node.

M is the number of output nodes.

$$\sigma(v_1) + \sigma(v_2) + \sigma(v_3) + \dots + \sigma(v_M) = 1$$

- Example:

$$v = \begin{bmatrix} 2 \\ 1 \\ 0.1 \end{bmatrix} \Rightarrow \sigma(v) = \begin{bmatrix} \frac{e^2}{e^2 + e^1 + e^{0.1}} \\ \frac{e^1}{e^2 + e^1 + e^{0.1}} \\ \frac{e^{0.1}}{e^2 + e^1 + e^{0.1}} \end{bmatrix} = \begin{bmatrix} 0.6590 \\ 0.2424 \\ 0.0986 \end{bmatrix}$$

# Derivative of Softmax function

$$y_i = \sigma(v_i) = \frac{e^{v_i}}{e^{v_1} + e^{v_2} + e^{v_3} + \dots + e^{v_M}} = \frac{e^{v_i}}{\sum_{k=1}^M e^{v_k}} = e^{v_i} \left( \sum_{k=1}^M e^{v_k} \right)^{-1}$$

$$\frac{\partial y_i}{\partial v_i} = \frac{\partial e^{v_i}}{\partial v_i} (\sum_{k=1}^M e^{v_k})^{-1} + e^{v_i} \frac{\partial (\sum_{k=1}^M e^{v_k})^{-1}}{\partial v_i},$$

$$= \frac{e^{v_i}}{\sum_{k=1}^M e^{v_k}} - \frac{e^{v_i} e^{v_i}}{(\sum_{k=1}^M e^{v_k})^2}$$

$$= \frac{e^{v_i}}{\sum_{k=1}^M e^{v_k}} \left( 1 - \frac{e^{v_i}}{\sum_{k=1}^M e^{v_k}} \right)$$

$$= \sigma(v_i) (1 - \sigma(v_i))$$

$$= y_i (1 - y_i)$$

$$\text{note } \frac{d(f(x)g(x))}{dx} = \frac{d(f(x))}{dx} g(x) + f(x) \frac{d(g(x))}{dx}$$

## Derivative of the **Binary** Cross-Entropy Loss


$$\begin{aligned}\frac{\partial L}{\partial \hat{y}} &= \frac{\partial (-y \ln(\hat{y}) - (1-y) \ln(1-\hat{y}))}{\partial \hat{y}} \\ &= -y \frac{1}{\hat{y}} - (1-y) \frac{-1}{1-\hat{y}} \\ &= \frac{-y(1-\hat{y}) + (1-y)\hat{y}}{\hat{y}(1-\hat{y})} \\ &= \frac{-(y-\hat{y})}{\hat{y}(1-\hat{y})}\end{aligned}$$

$$\hat{y} = \sigma(\hat{v})$$

$$\frac{\partial \hat{y}}{\partial \hat{v}} = \sigma(\hat{v})(1 - \sigma(\hat{v})) = \hat{y}(1 - \hat{y})$$

$$\frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial \hat{v}} = \frac{-(y-\hat{y})}{\hat{y}(1-\hat{y})} \times \hat{y}(1-\hat{y}) = -(y-\hat{y}) = \hat{y} - y = \sigma(\hat{v}) - y \rightarrow e$$

# Derivative of the Categorical Cross-Entropy Loss



One-Hot Encoding

$$\begin{aligned} & [y_1 \quad y_2 \quad y_3] \\ &= [1 \quad 0 \quad 0], \\ &= [0 \quad 1 \quad 0], \\ &= [0 \quad 0 \quad 1] \end{aligned}$$

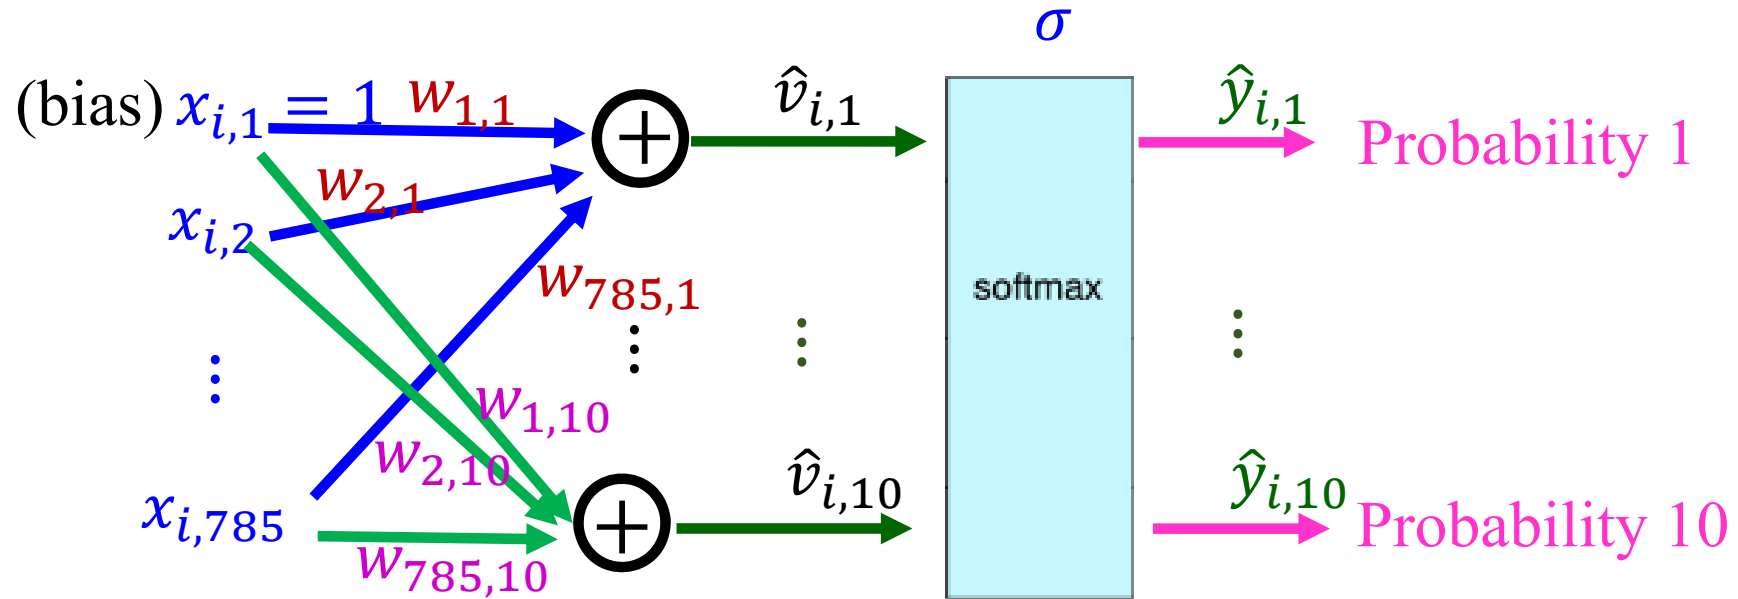
$$\frac{\partial L}{\partial v_j} = \frac{\partial}{\partial v_j} \left( -\sum_{k=1}^3 y_k \ln(\hat{y}_k) \right) = -\sum_{k=1}^3 y_k \frac{\partial \ln(\hat{y}_k)}{\partial v_j},$$

$$\hat{y}_k = \frac{e^{v_k}}{e^{v_1} + e^{v_2} + e^{v_3}}$$

$$\ln(\hat{y}_k) = v_k - \ln \left( \sum_{i=1}^3 e^{v_i} \right) \Rightarrow \frac{\partial \ln(\hat{y}_k)}{\partial v_j} = \frac{\partial v_k}{\partial v_j} - \frac{\partial}{\partial v_j} \ln \left( \sum_{i=1}^3 e^{v_i} \right) = 1\{k = j\} - \frac{e^{v_j}}{\sum_{i=1}^3 e^{v_i}}$$

$$\Rightarrow \frac{\partial L}{\partial v_j} = -\sum_{k=1}^3 y_k (1\{k = j\} - \hat{y}_j) = -\sum_{k=1}^3 y_k 1\{k = j\} + \hat{y}_j \sum_{k=1}^3 y_k = -y_j + \hat{y}_j \times 1 \equiv e_j$$

# Multiple inputs and outputs for Multiple Classification



$$\hat{v}_{i,j} = x_{i,1}w_{1,j} + x_{i,2}w_{2,j} + x_{i,3}w_{3,j} + \dots + x_{i,785}w_{785,j} = \mathbf{x}_i \mathbf{w}_j,$$

$$\text{where } \mathbf{x}_i = [x_{i,1} \ \dots \ x_{i,785}], \mathbf{w}_j = \begin{bmatrix} w_{1,j} \\ \vdots \\ w_{785,j} \end{bmatrix}, i = 1, \dots, N, j = 1, \dots, 10$$

$$\hat{y}_{i,j} = \sigma(v_{i,j}), \sigma \text{ is the softmax activation function}$$

**When  $b \equiv w_1 \neq 0$  and  $\hat{y} = \sigma(\hat{v})$  with CE loss**

$$\begin{bmatrix} 1 & x_{1,2} & \cdots & x_{1,785} \\ 1 & x_{2,2} & \cdots & x_{2,785} \\ \vdots & \vdots & \vdots & \vdots \\ 1 & x_{N,2} & \cdots & x_{N,785} \end{bmatrix} \begin{bmatrix} w_{1,1} & w_{1,2} & \cdots & w_{1,10} \\ w_{2,1} & w_{2,2} & \cdots & w_{2,10} \\ \vdots & \vdots & \vdots & \vdots \\ w_{785,1} & w_{785,2} & \cdots & w_{785,10} \end{bmatrix} = \begin{bmatrix} \hat{v}_{1,j} & \cdots & \hat{v}_{1,10} \\ \hat{v}_{2,j} & \cdots & \hat{v}_{2,10} \\ \vdots & \vdots & \vdots \\ \hat{v}_{N,j} & \cdots & \hat{v}_{N,10} \end{bmatrix}_{N \times 10}$$

$$\Rightarrow \hat{\mathbf{v}}_{N \times 10} = [\mathbf{1}_{N \times 1} \quad \mathbf{X}_{N \times 784}] \times \mathbf{w}_{785 \times 10}$$

$$\Rightarrow \hat{\mathbf{y}}_{N \times 10} = \sigma([\mathbf{1}_{N \times 1} \quad \mathbf{X}_{N \times 784}] \times \mathbf{w}_{785 \times 10})$$

$$\Rightarrow e_{1,j} = \sigma(\hat{v}_{1,j}) - y_{1,j} = \hat{y}_{1,j} - y_{1,j}, j = 1, \cdots 10$$

$$\vdots$$

$$\Rightarrow e_{N,j} = \sigma(\hat{v}_{N,j}) - y_{N,j} = \hat{y}_{N,j} - y_{N,j}, j = 1, \cdots 10$$

$$\Rightarrow \hat{\mathbf{e}}_{N \times 10} = \sigma([\mathbf{1}_{N \times 1} \quad \mathbf{X}_{N \times 784}] \times \mathbf{w}_{785 \times 10})$$

# Minimize Mean Cross Entropy Using Gradient Descent

- The loss function for output

$$\text{Loss}(\mathbf{w}) = \sum_{j=1}^{10} \left\{ \frac{-1}{N} \sum_{n=1}^N y_{n,j} \ln(\hat{y}_{n,j}) \right\}, \quad \frac{\partial \text{Loss}}{\partial \mathbf{w}_{i,j}} = \frac{\partial \text{Loss}}{\partial \hat{v}_{n,j}} \frac{\partial \hat{v}_{n,j}}{\partial \mathbf{w}_{i,j}}$$

Minimize  $\text{Loss}(\mathbf{w})$  w.r.t  $\mathbf{w}_{1,1}, \mathbf{w}_{2,1}, \dots, \mathbf{w}_{785,1}, \dots, \mathbf{w}_{1,10}, \mathbf{w}_{2,10}, \dots, \mathbf{w}_{785,10}$

$$\begin{aligned} \frac{\partial \text{Loss}}{\partial \mathbf{w}} &= \begin{bmatrix} \frac{\partial \text{Loss}}{\partial \mathbf{w}_{1,1}} & \frac{\partial \text{Loss}}{\partial \mathbf{w}_{1,2}} & \dots & \frac{\partial \text{Loss}}{\partial \mathbf{w}_{1,10}} \\ \frac{\partial \text{Loss}}{\partial \mathbf{w}_{2,1}} & \frac{\partial \text{Loss}}{\partial \mathbf{w}_{2,2}} & \dots & \frac{\partial \text{Loss}}{\partial \mathbf{w}_{2,10}} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial \text{Loss}}{\partial \mathbf{w}_{785,1}} & \frac{\partial \text{Loss}}{\partial \mathbf{w}_{785,2}} & \dots & \frac{\partial \text{Loss}}{\partial \mathbf{w}_{785,10}} \end{bmatrix} = \begin{bmatrix} \frac{1}{N} \sum_{n=1}^N e_{n,1} & \frac{1}{N} \sum_{n=1}^N e_{n,2} & \dots & \frac{1}{N} \sum_{n=1}^N e_{n,10} \\ \frac{1}{N} \sum_{n=1}^N e_{n,1} x_{n,2} & \frac{1}{N} \sum_{n=1}^N e_{n,2} x_{n,2} & \dots & \frac{1}{N} \sum_{n=1}^N e_{n,10} x_{n,2} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{1}{N} \sum_{n=1}^N e_{n,1} x_{n,785} & \frac{1}{N} \sum_{n=1}^N e_{n,2} x_{n,785} & \dots & \frac{1}{N} \sum_{n=1}^N e_{n,10} x_{n,785} \end{bmatrix} \\ &= \frac{1}{N} [\mathbf{1}_{N \times 1} \quad \mathbf{X}_{N \times 284}]_{785 \times N}^T \mathbf{e}_{N \times 10} \end{aligned}$$

- The **Batch Gradient Decent** method

$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} - \alpha \frac{\partial \text{Loss}}{\partial \mathbf{w}}$$



## Homework # 8-1, 8-2

- Re-write the Lecture 08\_4\_MNIST\_10\_classes\_sigmoid\_classification.jpynb in object-oriented format.

Hint: Homework 8\_1\_MINST\_10\_classes\_classification\_using\_logistic\_oop\_to\_do.jpynb

- Re-write the Lecture 08\_5\_MNIST\_10\_classes\_softmax\_classification.jpynb in object-oriented format.

Hint: Homework 8\_2\_MINST\_10\_classes\_classification\_using\_softmax\_oop\_to\_do.jpynb

**Deadline of Homework #8: 2022/11/21 3:30pm**